

# Import libraries

```
In [149... import pandas as pd
import numpy as np
import plotly.express as px
from sklearn.preprocessing import StandardScaler, MinMaxScaler
from sklearn.utils import resample
from sklearn.model_selection import train_test_split
from sklearn.neighbors import KNeighborsClassifier
from sklearn import metrics
from mlxtend.plotting import plot_confusion_matrix
from tqdm.notebook import tqdm
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
warnings.filterwarnings("ignore")
```

## Data loading and preprocessing

```
In [122... df = pd.read_csv("diabetes.csv")
```

```
In [123... df
```

```
Out[123...
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	Pediç
0	6	148	72	35	0	33.6	0
1	1	85	66	29	0	26.6	0
2	8	183	64	0	0	23.3	0
3	1	89	66	23	94	28.1	0
4	0	137	40	35	168	43.1	2
...	...	...	...	...	...	...	...
763	10	101	76	48	180	32.9	0
764	2	122	70	27	0	36.8	0
765	5	121	72	23	112	26.2	0
766	1	126	60	0	0	30.1	0
767	1	93	70	31	0	30.4	0

768 rows × 9 columns

```
In [124... df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 768 entries, 0 to 767
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Pregnancies           768 non-null   int64
1   Glucose               768 non-null   int64
2   BloodPressure         768 non-null   int64
3   SkinThickness         768 non-null   int64
4   Insulin               768 non-null   int64
5   BMI                   768 non-null   float64
6   Pedigree              768 non-null   float64
7   Age                   768 non-null   int64
8   Outcome               768 non-null   int64
dtypes: float64(2), int64(7)
memory usage: 54.1 KB

```

In [125... `df.describe().T`

```

Out[125...

```

	count	mean	std	min	25%	50%	
<b>Pregnancies</b>	768.0	3.845052	3.369578	0.000	1.00000	3.0000	6
<b>Glucose</b>	768.0	120.894531	31.972618	0.000	99.00000	117.0000	140
<b>BloodPressure</b>	768.0	69.105469	19.355807	0.000	62.00000	72.0000	80
<b>SkinThickness</b>	768.0	20.536458	15.952218	0.000	0.00000	23.0000	32
<b>Insulin</b>	768.0	79.799479	115.244002	0.000	0.00000	30.5000	127
<b>BMI</b>	768.0	31.992578	7.884160	0.000	27.30000	32.0000	36
<b>Pedigree</b>	768.0	0.471876	0.331329	0.078	0.24375	0.3725	0
<b>Age</b>	768.0	33.240885	11.760232	21.000	24.00000	29.0000	41
<b>Outcome</b>	768.0	0.348958	0.476951	0.000	0.00000	0.0000	1


In [126... `df["Outcome"].value_counts()`

```

Out[126...
0    500
1    268
Name: Outcome, dtype: int64

```

In [127... `sns.countplot(data=df, x=df["Outcome"])`  
`plt.show()`



## Upsampling

```
In [128... negative_data = df[df["Outcome"] == 0]
positive_data = df[df["Outcome"] == 1]
```

```
In [129... positive_upsample = resample(positive_data,
                                replace=True,
                                n_samples=int(0.9*len(negative_data)),
                                random_state=42)
```


```
In [130... new_df = negative_data
new_df = new_df.append(positive_upsample)
```

```
In [131... new_df.shape
```

```
Out[131... (950, 9)
```

```
In [132... new_df = new_df.sample(frac=1)
```

```
In [133... sns.countplot(data=new_df, x=new_df["Outcome"])
plt.show()
```



```
In [134... x = new_df.drop("Outcome", axis=1)
y = new_df[["Outcome"]]
```

```
In [135... scaler = MinMaxScaler()
scaled_values = scaler.fit_transform(x)
```

```
In [136... x_train, x_test, y_train, y_test = train_test_split(scaled_values, y, test_s
```

## KNN with elbow plot

```
In [137... k_values = [1, 3, 5, 7, 9, 11, 13, 15, 17, 19, 21, 23, 25, 27, 29, 31, 33, 35]
accuracy_values = []
```

```
In [138... for i in tqdm(range(len(k_values))):
    model = KNeighborsClassifier(n_neighbors=k_values[i])
    model.fit(x_train, y_train)
    y_pred = model.predict(x_test)
    accuracy = metrics.accuracy_score(y_test, y_pred)
    accuracy_values.append(accuracy)
```

```
0%|          | 0/25 [00:00<?, ?it/s]
```

```
In [143... px.line(x=k_values, y=accuracy_values)
```

```
In [144... optimal_k = -1
optimal_accuracy = -1
for i in list(zip(k_values, accuracy_values)):
    if i[1] > optimal_accuracy:
        optimal_k = i[0]
        optimal_accuracy = i[1]
```

```
In [145... knn_model = KNeighborsClassifier(n_neighbors=optimal_k)
```

```
In [146... knn_model.fit(x_train, y_train)
```


```
Out[146... ▼      KNeighborsClassifier
KNeighborsClassifier(n_neighbors=1)
```

```
In [147... y_pred = knn_model.predict(x_test)
```

```
In [148... print(metrics.classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.95	0.78	0.86	102
1	0.79	0.95	0.87	88
accuracy			0.86	190
macro avg	0.87	0.87	0.86	190
weighted avg	0.88	0.86	0.86	190

```
In [152... cm = metrics.confusion_matrix(y_test, y_pred)
plot_confusion_matrix(cm)
plt.show()
```



```
In [161... y_score = model.predict_proba(x_test)[: ,1]
```

```
In [162... false_positive_rate, true_positive_rate, threshold = metrics.roc_curve(y_test, y_score)
```

```
In [163... print('roc_auc_score for DecisionTree: ', metrics.roc_auc_score(y_test, y_score))
roc_auc_score for DecisionTree: 0.7575200534759358
```

```
In [165... plt.subplots(1, figsize=(10,7))
plt.title('Receiver Operating Characteristic - KNN')
plt.plot(false_positive_rate, true_positive_rate)
plt.plot([0, 1], ls="--")
plt.plot([0, 0], [1, 0], c=".7"), plt.plot([1, 1], c=".7")
plt.ylabel('True Positive Rate')
```

```
plt.xlabel('False Positive Rate')  
plt.show()
```

